METHODOLOGY, FORMALIZATION, AND NUMERICAL EVIDENCE (TECHNICAL / COMPUTATIONAL SECTION — DRAFT)

SHIN-YA KOYAMA AND SAAR SHAI

Abstract. Placeholder. The abstract for the full paper will be drafted jointly once the technical section is finalised.

1. Methodology, formalization, and numerical evidence

The section number `X` *and the equation / theorem / lemma tags (Lemma X.3.1, Theorem X.4.1, etc.) are placeholders to be assigned on integration into the full paper.*

This section records the technical and computational content of the paper: the algebraic identities at the heart of the corrected B_∞ framework (§ X.3–§ X.4), the open challenges that organise the program forward (§ X.4.4, § X.7), the numerical evidence at the two scales on which our verification operates (§ X.5), and the Lean 4 formalisation inventory (§ X.6).

The numerical evidence has **two distinct verified scales**, which we keep rigorously separate throughout. The **replication scale** is $x = 1.3 \cdot 10^{13}$, at which two independent prime-counting implementations agree on Koyama's Dominance-of--1 residue tables (§ X.5.1). The **analytic-identity scale** is $K \le 10^7$, at which the corrected B_∞ identity, the subleading constant C_1, and the Aoki–Koyama–Mertens drift toward $e^{-\gamma}$ are verified across three numerical stacks (§ X.5.2–§ X.5.4). Numbers from one scale are not transferred to the other.

1.1. Notation.

Fix a primitive non-principal Dirichlet character χ modulo $q \ge 2$, and let $\rho = \frac{1}{2} + i\tau$ with $\tau \ne 0$ be a simple non-trivial zero of $L(s, \chi)$. We use:

- $E_K(\chi, \rho) := \prod_{p \le K} (1 - \chi(p)\, p^{-\rho})^{-1}$ (truncated partial Euler product);
- $c_K(\chi, \rho) := \sum_{n \le K} \mu(n)\, \chi(n)\, n^{-\rho}$ (truncated Möbius sum);
- $D_K(\chi, \rho) := c_K(\chi, \rho)\, E_K(\chi, \rho)$ (Dirichlet D_K statistic);
- $T_K(\chi, \rho) := \sum_{p \le K} \sum_{k \ge 2} \chi(p)^k / (k\, p^{k\rho})$ and its limit T_∞ with $B_\infty := \exp(T_\infty)$;
- $C_1(\chi, \rho) := -L''(\rho, \chi)/(2\, L'(\rho, \chi)^2)$ (subleading constant);
- ψ for the primitive character of conductor $f \mid q$ inducing χ^2.

The branch of $\log L(2\rho, \psi)$ is fixed by analytic continuation from $\mathrm{Re}(s) > 1$, where the absolutely convergent log-Euler expansion applies, to the boundary line $\mathrm{Re}(s) = 1$ via Hadamard–de la Vallée Poussin non-vanishing.

1.2. Methodology of double verification.

Every numerical claim of § X.5 is computed by **two independent implementations in two languages with independent algorithms**.

- **L1 (primary).** `mpmath` 1.4 (Python 3.13), 50 decimal places. Direct partial-Möbius and partial-Euler evaluation; each zero ρ refined by Muller's method to $|L(\rho, \chi)| < 10^{-50}$.
- **L1b (in-language cross-check).** Same library, independent algorithm: Hurwitz-zeta expansion $L(s, \chi) = q^{-s} \sum_{a=1}^{q} \chi(a)\, \zeta(s, a/q)$ in place of `mpmath.dirichlet`, central-difference numerical derivatives at three step sizes, and an independent linear sieve for $\mu(n)$.
- **L2 (cross-language).** PARI/GP 2.17.3 (C), 57 dps default; with python-flint 0.8.0 / Arb (FLINT 3.3) at 250 bits as a third stack on the worst-case pair.

Date: May 13, 2026.

Acceptance gates: agreement to ≥ 12 significant digits on ρ and to ≥ 8 digits on L', L'', C_1; for residual quantities where cancellation occurs (the B_∞ residual), agreement to 10^{-12} on each of the four component sums separately. Branch choices and external citation provenance are recorded independently in L1 and L2.

For the Phase-1 prime-residue replication of § X.5.1, two analogous stacks **P1a** and **P1b** are used: `primesieve` 12.13 (segmented wheel sieve) and a hand-rolled plain-C segmented Eratosthenes sieve with no external dependency.

1.3. The local Perron double-pole residue.

Lemma X.3.1. *Let χ be a primitive non-principal Dirichlet character and let ρ be a simple zero of $L(s,\chi)$. Then for any $K > 1$,*

$$\operatorname*{Res}_{w=0}\left[\frac{K^w}{w\,L(w+\rho,\chi)}\right] \;=\; \frac{\log K}{L'(\rho,\chi)} \;+\; C_1(\chi,\rho). \tag{1}$$

Proof. Taylor-expand $L(w+\rho,\chi)$ at $w = 0$; invert to get $1/L(w+\rho,\chi) = (L'(\rho,\chi)\,w)^{-1} - L''(\rho,\chi)/(2\,L'(\rho,\chi)^2) + O(w)$. Multiply by $K^w/w = w^{-1} + \log K + O(w)$ and read off the coefficient of w^{-1}. □

Lemma X.3.1 is unconditional given simplicity of ρ. The full algebraic derivation is in Appendix B § B.2; the identity is also machine-verified in Lean 4 / Mathlib v4.28.0 (`LocalPerronResidue.lean`, 0 `sorry`; see § X.6).

1.4. Identities.

1.4.1. *X.4.1 Corrected B_∞ identity (unconditional).*

Our main algebraic result is the following unconditional identity for the prime-power tail T_∞.

Theorem X.4.1. *Let χ be a primitive non-principal Dirichlet character of conductor q, let ρ be a simple zero of $L(s,\chi)$ on the critical line, and let ψ be the primitive character of conductor $f \mid q$ inducing χ^2. Then*

$$T_\infty(\chi,\rho) \;=\; \tfrac{1}{2}\log L(2\rho,\psi) \;+\; \mathrm{BPC}_1(\chi,\rho) \;+\; \mathrm{BPC}_2(\chi,\rho) \;+\; T_{\geq 3}(\chi,\rho), \tag{2}$$

where

$$\mathrm{BPC}_1 = \tfrac{1}{2}\sum_{p\mid q,\, p\nmid f} \log(1 - \psi(p)\,p^{-2\rho}),$$

$$\mathrm{BPC}_2 = -\tfrac{1}{2}\sum_{k\geq 2}\frac{1}{k}\sum_p \frac{\chi(p)^{2k}}{p^{2k\rho}}, \qquad T_{\geq 3} = \sum_{k\geq 3}\frac{1}{k}\sum_p \frac{\chi(p)^k}{p^{k\rho}}.$$

The four right-hand-side terms are individually finite. The identity is unconditional given simplicity of ρ.

The $k = 1$ prime sum $\sum_p \chi^2(p)/p^{2\rho}$ is only conditionally convergent on $\mathrm{Re}(s) = 1$; its convergence is supplied by Akatsuka (2013, Lemma 2.1 / eq. (2.5)), which is itself unconditional (derived from PNT with explicit error term). BPC_2 and $T_{\geq 3}$ are absolutely convergent (minimum exponents $\mathrm{Re}(2k\rho) = 2$ and $\mathrm{Re}(k\rho) = \frac{3}{2}$). The full proof is given in Appendix A.

1.4.2. *X.4.2 Subleading constant C_1 and the partial Möbius identity.*

Theorem X.4.2. *Under the hypotheses of Theorem X.4.1, and assuming every off-target nontrivial zero ρ' of $L(s,\chi)$ with $|\gamma' - \tau| \leq T_K$ is simple (for a zero-avoiding height T_K as in the truncated explicit formula of Inoue 2021 Theorem 1),*

$$c_K(\chi,\rho) \;=\; \frac{\log K}{L'(\rho,\chi)} \;+\; C_1(\chi,\rho) \;+\; o(1) \qquad (K\to\infty). \tag{3}$$

The identity is unconditional in ρ given off-target simplicity. The rate $o(1) = O(K^{-1/2+\epsilon})$ is RH-conditional; the unconditional Soundararajan (2009) bound gives $o(1) = O(K^{-1/2}\exp((\log K)^{1/2}(\log\log K)^{14}))$.

The proof combines Inoue (2021, Theorem 1)'s truncated explicit formula for $M^*(K,\chi)$ with Lemma X.3.1 to extract the double-pole residue at ρ; off-target zero contributions are bounded by Soundararajan's argument. See Appendix B for the full proof.

1.4.3. *The Aoki–Koyama–Mertens constant (Hypothesis AK).* Aoki–Koyama (2023, *J. Number Theory* **245**, eq. (1.4), p. 235) states for a non-principal Dirichlet χ that, under DRH,

$$\lim_{x\to\infty}\Big((\log x)^m\prod_{p\le x}(1-\chi(p)/p^s)^{-1}\Big) \;=\; \frac{L^{(m)}(s,\chi)}{e^{m\gamma}\,m!}\cdot\begin{cases}\sqrt{2}, & \chi^2=1,\ s=\frac{1}{2},\\ 1, & \text{otherwise,}\end{cases}$$

with $m=m_\chi=\operatorname{ord}_{s=1/2}L(s,\chi)$. Specialised to the regime of this paper ($m=1$, $\rho\neq\frac{1}{2}$, branch multiplier 1), this reads:

$$\text{(AK)}\qquad \lim_{K\to\infty}E_K(\chi,\rho)\,\log K \;=\; \frac{L'(\rho,\chi)}{e^{\gamma}}.$$

This **corrects** the earlier target $L'(\rho,\chi)/\zeta(2)$: the ratio is $\zeta(2)/e^{\gamma}\approx 1.0828$. (AK) is DRH-conditional in the form stated by Aoki–Koyama.

1.4.4. *The conditional NDC limit (open).* Composing (3) with (AK) formally gives

$$\text{(NDC)}\qquad D_K(\chi,\rho)=c_K(\chi,\rho)\,E_K(\chi,\rho)\;\longrightarrow\; e^{-\gamma}\qquad (K\to\infty),$$

the corrected Numerical Duality Constant. The composition is mechanical provided (3) is strengthened to the **shifted Perron leading theorem**:

$$\text{(SP-L)}\qquad c_K(\chi,\rho)\;=\;\frac{\log K}{L'(\rho,\chi)}\;+\;o(\log K)\qquad (K\to\infty).$$

The obstruction is the off-target nontrivial-zero residue aggregate: even under DRH, an off-target zero λ of multiplicity $m\ge 2$ contributes

$$K^{\lambda-\rho}\,(\log K)^{m-1}/\big((m-1)!\,(\lambda-\rho)\,a_m(\lambda)\big),$$

and the $(\log K)^{m-1}$ factor is not absorbed by simplicity of ρ. A literature search (Inoue 2021, Soundararajan 2009, Ng 2004, Akatsuka 2013) did not surface a theorem that closes (SP-L); we state (NDC) as **conditional on (AK) and (SP-L)** (Question Q:Perron, § X.7).

1.5. Numerical findings.

1.5.1. *X.5.1 Phase-1 Dominance-of-−1 replication, $x=1.3\cdot 10^{13}$.* We independently replicate the prime-residue counts $\pi(x;N,a)$ of Koyama's *nontriv.pdf* for $N\in\{7,8,11,19,23\}$ and $x\in\{10^{12},1.3\cdot 10^{12},10^{13},1.3\cdot 10^{13}\}$ via two independent prime-enumeration implementations (`primesieve` 12.13 and a hand-rolled segmented C sieve). Headline numbers:

- $\pi(1.3\cdot 10^{13})=445{,}831{,}610{,}611$, cross-checked against `primesieve --count` standalone.
- Library-independence: at every one of the four checkpoints, the `primesieve` counts and the hand-rolled C-sieve counts agree on every residue class for every N.
- Hardware-independence: a second M1-class machine agrees through $x=1.3\cdot 10^{12}$ on every residue class for every N.
- Koyama's identity (3.1), a Dirichlet-orthogonality cross-check on the residue-count vector, is verified directly at all 495 (N,x,a)-cells (worst absolute residual $1.4\cdot 10^{-4}$).

Cell-by-cell comparison with Koyama's Tables 3–7 at all four checkpoints, all moduli:

Table	N	cells	exact	substantive disagreement (at $x = 1.3 \cdot 10^{13}$)
3	7	12	11	1 cell ($\Delta = 50$, clean digit-shift profile)
4	8	12	1	11 (small-x rows; possible x-label error in Table 4 draft)
5	11	20	19	1 cell ($a = 10$: our 11,503 vs Koyama 71,711)
6	19	18	15	3 (2 substantive at $a = 13, 18$; 1 sign flip at small x)
7	23	30	29	1 cell ($\Delta = 100$, clean transposition profile)
Total		**92**	**75**	**17** (74/81 91% excluding the 11 Table-4 small-x rows)

Comparing to the qualitative dominance-of--1 statement of Koyama (*nontriv.pdf*, § 3): the signal is **cleanly reproduced for** $N = 8$ ($-1 = 7$ is the strict largest at $1.3 \cdot 10^{13}$, $\pi(\cdot; 8, 7) - \pi(\cdot; 8, 1) =$ 164,958 vs 126,732 and 102,728). For $N = 19$ the ranking of -1 at this checkpoint is 3rd of 9 non-residues in both our run and Koyama's table — consistent with -1 being in the top group but not strictly dominant. For $N = 11$ the dominance turns on the disputed cell $a = 10$ (Table 5; see below); with Koyama's reported 71,711, -1 ranks 2nd of 5 (top group); with our 11,503, -1 ranks 3rd of 5 (outside top group). Pending reconciliation of that cell. For $N = 7$ and $N = 23$, Koyama himself notes (*nontriv.pdf* p. 19) that the predicted bias is not yet cleanly observed at this scale and attributes it to the exceptionally low-lying first zero of the relevant L-functions; our data agrees (-1 is 2nd of 3 for $N = 7$, mid-rank for $N = 23$). Koyama's illustrative estimate (*nontriv.pdf* p. 19, for $N = 19$) places the next sine-wave peak of the leading complex character at $x = e^{33.4} \approx 3.2 \cdot 10^{14}$, indicating that strict dominance for the harder moduli is expected only at substantially larger x.

The full replication bundle (source code, build hashes, TSV outputs, and reproducibility manifest) is deposited as Supplementary Material S1.

1.5.2. *Numerical values of the four Dirichlet pairs.* The four pairs (χ, ρ) used throughout § X.5.2–§ X.5.4 (all values computed in mpmath at 50 decimal places):

Pair	χ (conductor)	$\rho = \frac{1}{2} + i\tau$	$L'(\rho, \chi)$	$L''(\rho, \chi)$
χ_{-4}/z_1	χ_{-4} ($q = 4$)	$\frac{1}{2} + 6.020949i$	$1.296500 + 0.182765i$	$-1.697050 - 0.554017i$
χ_{-4}/z_2	χ_{-4}	$\frac{1}{2} + 10.243770i$	$1.788467 - 0.296776i$	$-3.319767 + 0.755548i$
χ_5	χ_5 ($q = 5$)	$\frac{1}{2} + 6.183578i$	$1.112930 - 0.448830i$	$-1.642973 + 1.035107i$

Pair	χ (conductor)	$\rho = \frac{1}{2} + i\tau$	$L'(\rho, \chi)$	$L''(\rho, \chi)$
χ_{11}	χ_{11} $(q = 11)$	$\frac{1}{2} + 3.547041i$	$1.696582 - 0.250988i$	$-3.121598 + 0.261219i$

L1b in-language cross-check (Hurwitz expansion, independent sieve) agrees with L1 to $|\Delta L'|, |\Delta L''| \lesssim 6 \cdot 10^{-12}$ and $|\Delta C_1| \lesssim 5 \cdot 10^{-13}$. L2 cross-language (PARI/GP 2.17.3) agrees with L1 to ≥ 11 decimal digits on every real and imaginary component. An Arb spot-check at 250 bits on the worst pair gives interval agreement on $|L'|$ within $3 \cdot 10^{-43}$.

1.5.3. *X.5.3 The Aoki–Koyama drift: $e^{-\gamma}$ vs $\zeta(2)^{-1}$. (Analytic-identity scale $K \leq 10^7$; not transferred from § X.5.1.)*
The modulus $|D_K|$ statistic at $K = 2 \cdot 10^6$ and $K = 10^7$ (40 dps, mean over the four pairs):

Quantity	$K = 2 \cdot 10^6$	$K = 10^7$	$\zeta(2)^{-1}$ target	$e^{-\gamma}$ target
Mean $|D_K| \cdot \zeta(2)$	0.992	0.974	1.000	$\zeta(2)\, e^{-\gamma} \approx 0.9237$
Mean $|E_K \log K|\, e^{\gamma}/|L'|$	n/a	0.942	n/a	1.000

The drift from 0.992 to 0.974 between $K = 2 \cdot 10^6$ and $K = 10^7$ is consistent with the AK normalisation $e^{-\gamma}$ at the natural $1/\log K$ finite-size scale and **incompatible with the $\zeta(2)^{-1}$ target** at the same scale. We do not claim convergence of the complex $D_K(\chi, \rho)$ from a modulus statistic alone — that depends on (SP-L), which is open.

1.5.4. *X.5.4 The B_∞ identity at the four pairs. (Analytic-identity scale $K \leq 10^7$; not transferred from § X.5.1.)*
Identity residual $|T_K - \mathrm{RHS}|$ for (2) at $K = 2 \cdot 10^6$ (mpmath, 50 dps) and $K = 10^7$ (PARI/GP 2.17.3, closed-form component evaluation):

Pair	$K = 2 \cdot 10^6$ (L1)	$K = 2 \cdot 10^6$ (L2)	$K = 10^7$ (L2)
χ_{-4}/z_1	$2.85 \cdot 10^{-3}$	$2.85 \cdot 10^{-3}$	$2.58 \cdot 10^{-3}$
χ_{-4}/z_2	$1.66 \cdot 10^{-3}$	$1.66 \cdot 10^{-3}$	$1.52 \cdot 10^{-3}$
χ_5	$4.24 \cdot 10^{-5}$	$4.24 \cdot 10^{-5}$	$1.22 \cdot 10^{-5}$
χ_{11}	$3.34 \cdot 10^{-5}$	$3.34 \cdot 10^{-5}$	$1.75 \cdot 10^{-5}$

L1 and L2 agree to all displayed digits at $K = 2 \cdot 10^6$ (stack difference $\leq 10^{-8}$). For the clean-character pairs χ_5, χ_{11} (where $\chi(2) \neq 0$ and there is no bad-prime contribution to BPC_1), the residual decays by a factor 1.9–3.5 from $K = 2 \cdot 10^6$ to $K = 10^7$, bracketing the $\sqrt{5} \approx 2.24$ predicted by the $K^{-1/2}/\log K$ rate of the boundary-line conditional tail (Akatsuka 2013 eq. (2.5)). The χ_{-4} pairs show systematically larger residuals consistent with the bad-prime $p = 2$ contribution to BPC_1.

1.5.5. *Two negative elliptic-curve findings.* We record two negative findings on the elliptic-curve side, both deferred to the supplementary record for the design data, source, and resampling diagnostics:

- *Conductor-confounded rank trend.* A multivariate OLS regression of $E[C_1^2]$ on $(\mathrm{rank}, \log N)$ across 19 weight-2 elliptic curves shows a stable $\log N$ coefficient and an *unstable* rank coefficient (bootstrap 95% CI for the rank coefficient includes zero; one rank-3 anchor swings the coefficient by 63%). We describe this as a conductor-confounded trend, not a rank law (Question Q:conductor, § X.7).

- *Raw* $\mathrm{Sym}^2/\langle f, f\rangle$ *proportionality.* The raw ratio ranges over seven orders of magnitude across $\{37a_1, 389a_1, \Delta\}$ and is empirically falsified in its raw form (Question Q:Sym2, § X.7).

1.6. **Lean 4 / Mathlib4 formalisation.** We accompany the paper with a Lean 4 / Mathlib4 lake project (`formal-conjectures/` directory of the supplementary archive). The toolchain is `leanprover/lean4:v4.28.0`; Mathlib is pinned at commit `8f9d9cff6bd728b17a24e163c9402775d9e6a365`. The Lean inventory fixes the **statements** of every identity of § X.4, ensures normalisations and branch conventions are syntactically explicit, and records each statement's proof status against a public audit trail.

Build status. `lake build FormalConjectures` succeeds on all **9 files** in `formal-conjectures/` with **5 `sorry` warnings**, each annotated in-source as `MATHLIB-PREREQ:` or `RESEARCH-OPEN:`. **Six files are fully proved (0 `sorry`)**: all the algebraic content of § X.3, § X.4, the smoothed explicit-formula chain, the Mertens spectroscope universality statement, the Farey bridge identity, and DPAC for $K \leq 4$. No `axiom` is introduced anywhere in the project. The full per-`sorry` inventory is in `LEAN_SORRY_STATUS.md` of the reproducibility bundle.

Status of headline theorems in § X.3–§ X.4:

Paper object	Lean status
Lemma X.3.1 (local Perron residue)	**Lean-verified unconditional** (`LocalPerronResidue.lean`, 0 `sorry`)
Theorem X.4.1 (corrected B_∞ identity)	**Lean-verified conditional** on the convergence-of-partial-sum hypothesis derived in Appendix A (`CorrectedBInfty.lean`, 0 `sorry`)
Theorem X.4.2 (c_K leading + subleading)	Pen-and-paper proof in Appendix B; off-target-zero-simplicity hypothesis stated explicitly
(AK), (SP-L), (NDC)	(AK) cited (Aoki–Koyama 2023); (SP-L) and (NDC) open challenges (Q:Perron, § X.7)
DPAC	**Lean-verified for $K \leq 4$ unconditional** (`DPAC_closure_attempt.lean`, 0 `sorry`); general K open (LI-class), with obstruction certificate via Pólya 1913

Paper object	Lean file	Status
Boundary residue $R_0 = -2$ for a Gaussian-cutoff Mellin-shift explicit formula (companion strand) and its algebraic-glue chain	`SmoothedDwfFormula_full.lean`	**THEOREM (chain), 0 `sorry`.** All 17 algebraic-glue lemmas closed unconditionally; the two analytic prerequisites `mellin_decay` (Stirling on Γ vertical strips) and `inv_zeta_polynomial_growth` (Titchmarsh § 3.11) are now stated as explicit hypotheses on the theorems that consume them, both Mathlib v4.28.0 gaps.
Lemma X.3.1 (local Perron residue)	`LocalPerronResidue.lean`	**THEOREM (0 `sorry`).** The residue identity is stated as a `Tendsto` limit at L analytic with simple zero at 0 (the general-ρ case reduces by $L \mapsto L(\cdot + \rho)$).

Paper object	Lean file	Status
Theorem X.4.1 (B_∞ identity)	`CorrectedBInfty.lean`	**THEOREM (0 `sorry`), conditional on `h_convergence`.** The four-component identity is proved against `noncomputable defs` of T_∞, $T_{\geq 3}$, BPC_1, BPC_2, L given an added hypothesis `h_convergence : Tendsto T_K atTop (nhds (RHS))`. This hypothesis packages exactly the four analytic inputs of the pen-and-paper proof in Appendix A (Akatsuka 2013, log-Euler-product, imprimitive induction, geometric tails); the Lean proof uses `Classical.epsilon_spec` + `tendsto_nhds_unique` and is three lines.
Farey bridge identity	`FareyBridgeIdentity.lean`	**THEOREM (0 `sorry`), conditional on a Ramanujan-sum decomposition hypothesis** packaging the Hardy–Wright Theorem 304 input (`c_q(p) = \mu(q)` for $\gcd(p, q) = 1$). Two helper lemmas (`mertens_eq_pred_add_moebius`, `sum_moebius_Icc_eq_mertens`) closed without `sorry`. `MATHLIB-PREREQ: Mathlib.NumberTheory.RamanujanSum` for the unconditional version.
Mertens spectroscope universality	`MertensSpectroscopeUniversality.lean`	**THEOREM (0 `sorry`), conditional on an explicit-formula asymptotic hypothesis** (Soundararajan 2009 Theorem 1 input). Proof: `Filter.tendsto_atTop.mpr h_explicit_formula`.
Farey sign pattern	`FareySignPattern.lean`	**NEGATIVE.** The pointwise version is falsified at $p = 237{,}733$ and $p = 243{,}799$ (recorded as theorems, not axioms — the project's "no `axiom`" convention is preserved). The density-one surviving version is stated as a `Tendsto`. 3 `sorry`s.

Paper object	Lean file	Status
Dirichlet Polynomial Avoidance (DPAC) — partial closure + bridges	`DPAC_closure_attempt.lean`, `DirichletPolynomialAvoidance.lean`, `DPAC_full.lean`	**PARTIAL.** `DPAC_closure_attempt.lean` (0 `sorry`) proves DPAC unconditionally for $K \in \{2, 3, 4\}$ using only $0 < \mathrm{Re}(\rho) < 1$ (`dpac_K_eq_2`, `dpac_K_eq_3`, `dpac_K_eq_4`, `dpac_le_4`). It also reformulates the open case as `FiniteLogRatioLI` and records the obstruction certificate (Pólya 1913 discreteness of the exponential-polynomial zero set + a single open avoidance statement). The headline conjecture for general K remains `sorry` in `DPAC_full.lean:297` and `DirichletPolynomialAvoidance.lean:48`, diagnostically comparable to the Linear Independence Hypothesis for ζ-zero ordinates; the four explicit phase-avoidance bridges (`dpac_of_logPrimePhaseAvoidance`, . . . , `dpac_of_certifiedZetaZeroSample`) are closed without `sorry`.

The role of the Lean artifact is to fix the statements and provide a publicly inspectable audit trail of the proof obligations remaining.

1.7. **Open challenges.** The following structure the next phase of the program.

Q:Perron (Shifted Perron leading theorem). Prove (SP-L) ((SP-L)) for primitive non-principal χ and simple non-central ρ.

Sufficient packages: (a) all off-target zeros simple, and $Z_{\mathrm{simple}}(K, T_K) := \sum_{\rho' \neq \rho, |\gamma'| \leq T_K} K^{\rho' - \rho}/[(\rho' - \rho)\, L'(\rho', \chi)] = o(\log K)$ at a zero-avoiding height $T_K \asymp K(\log K)^{-B}$; (b) a Dirichlet shifted second moment $\sum_{\rho'}^{\mathrm{mult}} |L(\rho' + \alpha, \chi)|^{-2} \ll_\chi (\log K)^{O(1)}$. The total-Möbius bounds of Soundararajan type are too coarse to isolate the pointwise cancellation at the $\log K$ scale.

A naïve GL(1) transfer of the GL(2) halo bound (cluster-summed contour residues over a halo region, giving $\ll T^{7/4+\varepsilon}$ at T-scale) yields only $K^{1/2+\varepsilon}$ for the off-target residue aggregate at K-scale, far above the $o(\log K)$ target — the halo route in its present form does not close (SP-L). The cluster-summed-residue pivot does, however, replace the rooted Palm wall that obstructs the termwise estimate. Details in the supplementary record.

Q:EC-recip (GL(2) reciprocal-derivative control). Prove a fixed-curve theorem for $\sum_\gamma \widehat{W}(i\gamma)\, e^{i\gamma u}/L'(E, 1 + i\gamma)$ giving cancellation $o(u^r)$, or a minimum-modulus estimate on a vertical line with explicit exponent < 2.

Without a GL(2) analogue of Aoki–Koyama, the EC side remains at the level of quantitative ensemble evidence.

Q:DPAC (Dirichlet Polynomial Avoidance). Prove DPAC for general K. We give an unconditional Lean proof for $K \in \{2, 3, 4\}$; the general case reduces, via

Pólya 1913 discreteness of the finite-exponential-polynomial zero set, to a single open avoidance statement at ζ-zero ordinates diagnostically comparable to the Linear Independence Hypothesis.

Further questions (from the EC and ensemble negatives of § X.5.5; deferred to the supplementary record):

- *Q:conductor* — Replicate (the rank-vs-log N regression of § X.5.5) on a curve set in which rank and $\log N$ are not collinear, to separate the conductor contribution from any genuine rank dependence.
- *Q:Sym2* — Identify a completed / archimedean-corrected Sym^2 normalisation replacing the empirically falsified raw $\mathrm{Sym}^2/\langle f, f\rangle$ proportionality.
- *Q:EC-NDC* — Find a normalisation of D_K^E for which the universal limit exists and survives a null-control gate. The sharp-cutoff form $D_K^E \cdot \zeta(2) \to 1$ is falsified through $K = 10^6$; smoothed variants pass empirically but also pass a null-control gate against predeclared null transformations.

1.8. **Code, data, and certificate availability.** All scripts, refined zero data, numerical-table CSVs, convergence logs, the Lean 4 lake project, and the reproducibility manifest will be deposited as a single self-contained reproducibility bundle (Supplementary material S1), mirrored at a Zenodo DOI at acceptance. The bundle pins all software versions: Lean toolchain `leanprover/lean4:v4.28.0`, Mathlib commit `8f9d9cff6bd728b17a24e163c9402775d9e6a365`, `mpmath` 1.4, PARI/GP 2.17.3, FLINT 3.3 / python-flint 0.8.0. The Phase-1 Dominance-of--1 replication bundle is included as the Supplementary Material S1 archive. Each numerical table in § X.5 cites the L1 script and L2 reproducer; each external theorem cited in § X.4 has its PDF retrieval recipe, page/equation, and verbatim quote recorded in the citation audit (Supplementary S2, *Citation audit*).

1.9. **References (section).** External references cited in § X.3–§ X.7. Full page-and-equation provenance for each is in Supplementary S2 (citation audit).

- **Akatsuka, H.** (2013). *The Euler product for the Riemann zeta function in the critical strip.* The boundary-line Mertens-type partial-summation estimate $\sum_{p\le X} \chi(p)/p^{1+2i\tau} = c(\chi,\tau) + O(1/\log X)$ used in § X.4.1, Appendix A § A.2.3; Lemma 2.1 / eq. (2.5). Unconditional (derived from PNT with explicit error term).
- **Aoki, T. and Koyama, S.** (2023). *Generalised Mertens constant for Dirichlet L-functions.* J. Number Theory **245**, 233–262; eq. (1.4), p. 235. The $e^{-m\gamma}$-normalised partial-Euler-product asymptotic at noncentral zeros under DRH. Cited as Hypothesis (AK) in § X.4.3.
- **Davenport, H.** *Multiplicative Number Theory*, 3rd ed., GTM **74**, Springer 2000. Hadamard–de la Vallée Poussin non-vanishing on $\mathrm{Re}(s) = 1$.
- **Hardy, G.H. and Wright, E.M.** *An Introduction to the Theory of Numbers*, 6th ed., Oxford 2008. Theorem 304 (Ramanujan-sum identity at primes).
- **Inoue, S.** (2021). *Truncated explicit formula for $M(x, \chi)$.* Journal de Théorie des Nombres de Bordeaux **33**, Theorem 1. Used in Appendix B § B.1, § B.3 to set up the contour integration for Theorem X.4.2.
- **Montgomery, H.L. and Vaughan, R.C.** *Multiplicative Number Theory I. Classical Theory*, Cambridge 2007. Theorem 9.4 (textbook bad-prime correction structure).
- **Ng, N.** (2004). *The distribution of the summatory function of the Möbius function.* Proc. London Math. Soc. **89**(3), 361–389. Discussed alongside Soundararajan 2009 for the off-target zero aggregate (Question Q:Perron, § X.7).
- **Pólya, G.** (1913). *Über die Nullstellen gewisser ganzer Funktionen.* Classical discreteness of the zero set of a finite exponential polynomial; used in `DPAC_closure_attempt.lean` for the obstruction certificate.

- **Soundararajan, K.** (2009). *Partial sums of the Möbius function.* Ann. of Math. **170**(2), 1409–1422; Theorem 1. The RH-conditional and the unconditional rate bounds for the Möbius partial sum used in Theorem X.4.2 and Appendix B § B.4.
- **Tenenbaum, G.** *Introduction to Analytic and Probabilistic Number Theory*, 3rd ed., GSM **163**, AMS 2015. Chapter II.5 (Hadamard–de la Vallée Poussin treatment).
- **Titchmarsh, E.C.** *The Theory of the Riemann Zeta-Function*, 2nd ed., Heath-Brown rev., OUP 1986. § 3.11 (the polynomial-growth bound on $1/\zeta(s)$ used in `SmoothedDwfFormula_full.lean` and cited as a Mathlib prerequisite in § X.6).

Appendix A. Appendix A — Full proof of Theorem X.4.1 (B_∞ identity)

This appendix gives the full proof of Theorem X.4.1 of the main text: for a primitive non-principal Dirichlet character χ of conductor q and a simple zero ρ of $L(s,\chi)$ on the critical line, with ψ denoting the primitive character of conductor $f \mid q$ that induces χ^2,

$$T_\infty(\chi,\rho) \;=\; \tfrac{1}{2}\log L(2\rho,\psi) \;+\; \mathrm{BPC}_1(\chi,\rho) \;+\; \mathrm{BPC}_2(\chi,\rho) \;+\; T_{\geq 3}(\chi,\rho), \tag{$\star$}$$

with the four terms on the right defined in § X.1 of the main text. The proof uses no hypothesis beyond simplicity of ρ as a zero of $L(s,\chi)$; in particular it does not use RH, GRH, EDRH, or DRH.

A.1. **Setup — partial-sum decomposition.** For a primitive non-principal Dirichlet character χ modulo $q \geq 2$ and a simple zero $\rho = \frac{1}{2} + i\tau$ with $\tau \neq 0$ of $L(s,\chi)$ on the critical line, define

$$T_K(\chi,\rho) \;:=\; \sum_{p \leq K} \sum_{k \geq 2} \frac{\chi(p)^k}{k\, p^{k\rho}}.$$

Inside the partial sum, the order of summation may be swapped unconditionally because for each fixed prime $p \geq 2$,

$$\sum_{k \geq 2} \frac{(\chi(p)\, p^{-\rho})^k}{k} \;=\; -\log\bigl(1 - \chi(p)\, p^{-\rho}\bigr) \;-\; \chi(p)\, p^{-\rho},$$

with $\bigl|\chi(p)\, p^{-\rho}\bigr| \;=\; p^{-1/2} \;\leq\; 2^{-1/2} \;<\; 1$, so the inner series is absolutely convergent and the logarithm is taken along the principal branch with no ambiguity. Therefore we may re-index the partial sum by k first:

$$T_K(\chi,\rho) \;=\; \tfrac{1}{2} \sum_{p \leq K} \frac{\chi(p)^2}{p^{2\rho}} \;+\; \sum_{k \geq 3} \frac{1}{k} \sum_{p \leq K} \frac{\chi(p)^k}{p^{k\rho}}. \tag{4}$$

The only delicate object in passing to $K \to \infty$ is the $k = 2$ piece, since $\mathrm{Re}(2\rho) = 1$ lies on the boundary line where the Dirichlet prime sum $\sum_p \chi^2(p) p^{-s}$ is conditionally (not absolutely) convergent. We address this in § A.2. The $k \geq 3$ piece is absolutely convergent and is handled in § A.3.

A.2. A.2 Identification of the $k = 2$ partial sum.

A.2.1. *A.2.1 The squared character χ^2 and its primitive companion ψ.* For χ primitive of conductor q, define χ^2 by $\chi^2(n) := \chi(n)\,\chi(n)$ for $\gcd(n,q) = 1$ and $\chi^2(n) = 0$ otherwise. The character χ^2 is a (possibly imprimitive) Dirichlet character modulo q. By a standard textbook result (e.g. Montgomery–Vaughan, *Multiplicative Number Theory I*, Theorem 9.4 / Corollary 9.5), every Dirichlet character modulo q is induced from a unique primitive character ψ of some conductor $f \mid q$, in the sense that $\chi^2(n) = \psi(n)$ for $\gcd(n,q) = 1$ and $\chi^2(n) = 0$ for $\gcd(n,q) > 1$. The corresponding L-function identity is

$$L(s,\chi^2) \;=\; L(s,\psi) \cdot \prod_{p \mid q,\; p \nmid f} \Bigl(1 - \frac{\psi(p)}{p^s}\Bigr), \tag{5}$$

valid initially for $\mathrm{Re}(s) > 1$ and then by analytic continuation in the natural domain of the right-hand side.

A.2.2. *A.2.2 Log-Euler-product expansion of* $\log L(s, \chi^2)$. For $\mathrm{Re}(s) > 1$, taking the principal branch of the logarithm of the Euler product $L(s, \chi^2) = \prod_{p \nmid q}(1 - \chi^2(p)p^{-s})^{-1}$ gives

$$\log L(s, \chi^2) \;=\; \sum_{p \nmid q} \sum_{k \geq 1} \frac{\chi^2(p)^k}{k\, p^{ks}} \;=\; \sum_{p} \sum_{k \geq 1} \frac{\chi(p)^{2k}}{k\, p^{ks}}, \tag{6}$$

where the second equality extends the prime sum to *all* primes (the contribution of $p \mid q$ vanishes because $\chi(p) = 0$ for those primes). The double sum is absolutely convergent in $\mathrm{Re}(s) > 1$.

Isolating the $k = 1$ term in (6) yields

$$\sum_{p} \frac{\chi(p)^2}{p^s} \;=\; \log L(s, \chi^2) \;-\; \sum_{k \geq 2} \frac{1}{k} \sum_{p} \frac{\chi(p)^{2k}}{p^{ks}}, \qquad \mathrm{Re}(s) > 1. \tag{7}$$

A.2.3. *A.2.3 Analytic continuation of (7) to* $s = 2\rho$. We extend (7) to the half-plane $\mathrm{Re}(s) > \frac{1}{2}$ minus the singular set of the right-hand side, as follows.

Right-hand side analyticity at $s = 2\rho$. The function $\log L(s, \chi^2)$ is holomorphic at $s = 2\rho$ provided $L(2\rho, \chi^2) \neq 0$ and $L(s, \chi^2)$ has no pole there. When χ^2 is principal (i.e., $f = 1$, hence χ^2 is induced from the trivial character), $L(s, \psi) = \zeta(s)$, which has a simple pole at $s = 1$. But $s = 2\rho = 1 + 2i\tau$ with $\tau \neq 0$, so $s \neq 1$ and the pole of ζ is avoided. Non-vanishing of $L(s, \chi^2)$ on the line $\mathrm{Re}(s) = 1$ (away from the potential pole at $s = 1$) is the classical Hadamard–de la Vallée Poussin theorem (see e.g. Tenenbaum, *Introduction to Analytic and Probabilistic Number Theory*, Chapter II.5). Hence $\log L(s, \chi^2)$ admits an analytic continuation through the neighbourhood of $s = 2\rho$ along the principal branch.

Right-hand $k \geq 2$ sum analyticity. For $k \geq 2$, $\mathrm{Re}(ks) \geq k \geq 2$ on the line $\mathrm{Re}(s) = 1$, so the sum $\sum_{k \geq 2}(1/k) \sum_p \chi(p)^{2k} p^{-ks}$ converges absolutely on $\mathrm{Re}(s) > 1/2$ and defines a holomorphic function there.

Left-hand side analyticity. The prime sum $\sum_p \chi^2(p)p^{-s}$ on the line $\mathrm{Re}(s) = 1$ (excluding $s = 1$ if χ^2 is principal) converges conditionally by a partial-summation argument due to Akatsuka (2013), *The Euler product for the Riemann zeta function in the critical strip*, Lemma 2.1 and equation (2.5), which establishes that for $t_0 \neq 0$,

$$\sum_{p \leq X} \frac{1}{p^{1+2it_0}} \;=\; c(t_0) \;+\; O\big((\log X)^{-1}\big), \tag{8}$$

proved by partial summation against the prime number theorem with explicit error term. The estimate (8) is **unconditional** (it does not require RH or any GRH-type hypothesis); the same partial-summation argument applies with $\chi^2(p)/p^{1+2i\tau}$ in place of p^{-1-2it_0}, with the character orthogonality producing the appropriate cancellation. We denote the limit by $\Sigma_2(\chi, \rho) := \lim_{X \to \infty} \sum_{p \leq X} \chi^2(p)\, p^{-2\rho}$.

We pass to the boundary $s = 2\rho$ by Abel summation. For $\sigma > 1$, identity (7) reads $\log L(s, \chi^2) - \sum_{k \geq 2}(1/k) \sum_p \chi(p)^{2k}/p^{ks} = \sum_p \chi(p)^2/p^s$, the left side analytic in a neighbourhood of $s = 2\rho$ (here we use Hadamard–de la Vallée Poussin: $L(s, \psi) \neq 0$ on $\mathrm{Re}(s) = 1$ with $s \neq 1$, where ψ is the primitive character inducing χ^2, together with the imprimitive Euler-factor identity $L(s, \chi^2) = L(s, \psi) \prod_{p \mid q,\, p \nmid f}(1 - \psi(p)/p^s)$). Writing $S(X) := \sum_{p \leq X} \chi(p)^2/p^{2\rho}$ and applying Abel summation gives, for $\sigma > 1$, $\sum_p \chi(p)^2/p^s = (s - 2\rho) \int_2^\infty S(X)\, X^{-(s-2\rho)-1}\, dX + \lim_{X \to \infty} S(X)\, X^{-(s-2\rho)}$; Akatsuka (2013) Lemma 2.1 / eq. (2.5) provides $S(X) = c(\chi, \rho) + O(1/\log X)$ unconditionally as $X \to \infty$, so the limit term equals $c(\chi, \rho)$ at $s = 2\rho$ and the integral converges to a continuous function of s in a neighbourhood. Hence (7) extends by continuity to $s = 2\rho$ as the boundary-value identity

$$\Sigma_2(\chi,\rho) \;=\; \log L(2\rho,\chi^2) \;-\; \sum_{k\ge 2}\frac{1}{k}\sum_p \frac{\chi(p)^{2k}}{p^{2k\rho}}. \tag{9}$$

A.2.4. *A.2.4 Multiplying by* $1/2$. Multiplying (9) by $\frac{1}{2}$ and rewriting $\log L(2\rho,\chi^2)$ via the imprimitive-induction identity (5) gives

$$\tfrac{1}{2}\Sigma_2(\chi,\rho) \;=\; \tfrac{1}{2}\log L(2\rho,\psi) \;+\; \tfrac{1}{2}\sum_{p\mid q,\ p\nmid f}\log\bigl(1-\psi(p)\,p^{-2\rho}\bigr) \;-\; \tfrac{1}{2}\sum_{k\ge 2}\frac{1}{k}\sum_p\frac{\chi(p)^{2k}}{p^{2k\rho}}. \tag{10}$$

The bad-prime correction in the middle term is precisely the definition of BPC_1 from § X.1; the last term is BPC_2.

A.3. **A.3 Absolute convergence of the $k \ge 3$ tail.** For each fixed prime p and each $k\ge 3$, $|\chi(p)^k p^{-k\rho}| \le p^{-k/2}$, so the inner k-series is bounded by a geometric series:

$$\Bigl|\sum_{k\ge 3}\frac{\chi(p)^k}{k\,p^{k\rho}}\Bigr| \;\le\; \sum_{k\ge 3}\frac{p^{-k/2}}{k} \;\le\; \frac{p^{-3/2}}{3}\cdot\frac{1}{1-p^{-1/2}}.$$

Summing over primes:

$$|T_{\ge 3}| \;\le\; \frac{1}{3}\sum_p\frac{p^{-3/2}}{1-p^{-1/2}} \;\le\; \frac{1}{3(1-2^{-1/2})}\sum_p p^{-3/2} \;=\; \frac{P(3/2)}{3(1-2^{-1/2})} \;\approx\; 0.515,$$

where $P(3/2)=\sum_p p^{-3/2}\approx 0.45224$ is the prime zeta function at $s=3/2$. The truncation tail is also estimable: by partial summation against PNT, for $\alpha>1$, $\sum_{p>K}p^{-\alpha}\le\frac{2K^{1-\alpha}}{(\alpha-1)\log K}$, so

$$|T_{\ge 3}-T_{\ge 3,K}| \;\le\; \frac{2}{3(1-2^{-1/2})}\cdot\frac{K^{-1/2}}{\log K},$$

which is $\sim 3\cdot 10^{-4}$ at $K=10^6$ and $\sim 1\cdot 10^{-4}$ at $K=2\cdot 10^6$.

A.4. **Assembling the identity.** Define

$$T_\infty(\chi,\rho) \;:=\; \lim_{K\to\infty} T_K(\chi,\rho).$$

The limit exists by combining the two parts:

- The $k=2$ part is $\frac{1}{2}\Sigma_2(\chi,\rho)$, whose limit exists by Akatsuka (2013) Lemma 2.1 / eq. (2.5) and equals the right-hand side of (10).
- The $k\ge 3$ part is absolutely convergent by § A.3.

Hence

$$T_\infty(\chi,\rho) \;=\; \tfrac{1}{2}\log L(2\rho,\psi) \;+\; \mathrm{BPC}_1(\chi,\rho) \;+\; \mathrm{BPC}_2(\chi,\rho) \;+\; T_{\ge 3}(\chi,\rho),$$

which is identity $(\star)$. □

A.5. **Bad-prime correction for the four computational pairs.** For the four (χ,ρ) pairs used in the numerical work of § X.5.4, the bad-prime correction BPC_1 takes the following character-explicit forms:

Character	q	χ^2 order	primitive ψ inducing χ^2	f	bad primes $\{p \mid q,\ p\nmid f\}$
χ_{-4}	4	1 (principal mod 4)	trivial character ($\Rightarrow L=\zeta$)	1	$\{2\}$
χ_5	5	2	Legendre symbol $(\cdot/5)$	5	$\emptyset$

Character	q	χ^2 order	primitive ψ inducing χ^2	f	bad primes $\{p \mid q,\ p \nmid f\}$
χ_{11}	11	5	order-5 character mod 11	11	$\emptyset$

Explicitly: - For χ_{-4}: $\mathrm{BPC}_1 = \frac{1}{2}\log(1 - 2^{-2\rho})$. For $\rho = \frac{1}{2} + i\tau$, this is $\frac{1}{2}\log(1 - 2^{-1-2i\tau})$, of modulus $\le \frac{1}{2}\log(1/(1-1/2)) = \frac{1}{2}\log 2 \approx 0.347$. - For χ_5 and χ_{11}: $\mathrm{BPC}_1 = 0$.

This explains the numerical observation in § X.5.4 that the χ_{-4} pairs show slower convergence of the B_∞ identity residual than χ_5 and χ_{11}: the bad-prime contribution from $p = 2$ adds an extra layer of conditionally-convergent boundary mass to the χ_{-4} identity that is absent for χ_5, χ_{11} (the latter pure conditional-tail decay scales as $K^{-1/2}/\log K$ exactly).

A.6. **A.6 What the proof does *not* use.** For clarity:

- The proof does **not** use RH, GRH, EDRH, or DRH. The only conditional convergence on the boundary line $\mathrm{Re}(s) = 1$ is the $k = 1$ prime sum $\sum_p \chi^2(p)p^{-2\rho}$, handled by Akatsuka 2013 Lemma 2.1 / eq. (2.5), which itself is *unconditional* — derived from PNT with explicit error term.
- The proof does **not** use any rate of convergence for the partial Möbius / spectroscope sum $c_K(\chi, \rho)$. The B_∞ identity is a statement about the limit T_∞; it makes no claim about the *speed* at which $T_K \to T_\infty$.
- The proof does **not** establish the Aoki–Koyama–Mertens limit ((AK)) for $E_K \log K$, which is a separate statement on the *additive* side and is cited as Hypothesis AK (with verbatim quotation in § X.4.3).
- The proof does **not** establish the conditional NDC limit ((NDC)); for that, Hypothesis AK and the shifted Perron leading hypothesis (SP-L) are both needed (see § X.4.4 and § X.7).

A.7. **Lean formalisation.** The four-component identity is formalised in Lean 4 / Mathlib v4.28.0 as the theorem `corrected_B_infty` in `formal-conjectures/CorrectedBInfty.lean` of the supplementary archive. The Lean proof is parameterised by an explicit `Filter.Tendsto` hypothesis asserting that the partial-sum sequence $T_K(\chi, \rho)$ converges to the four-component right-hand side as $K \to \infty$. **The convergence-as-hypothesis is precisely the conclusion of the present appendix** (7) + (9) + the $k \ge 3$ tail of § A.3. Given the convergence, the Lean proof is three lines: `unfold T_inf` (the `Classical.epsilon` of the `Tendsto` predicate), `Classical.epsilon_spec` (which yields that T_∞ inherits the same limit), and `tendsto_nhds_unique` (since $\mathbb{C}$ is Hausdorff). The file uses only the standard `propext`, `Classical.choice`, `Quot.sound` axioms.

A fully unconditional Lean proof of `corrected_B_infty` (i.e., one that *derives* the convergence rather than taking it as a hypothesis) requires upstream Mathlib formalisations of Akatsuka 2013 eq. (2.5) and the imprimitive-induction Euler-factor identity for $L(s, \chi^2)$; both are `MATHLIB-PREREQ` and not yet upstream as of v4.28.0.

A.8. **Numerical verification (summary; full table in § X.5.4).** The identity ($\star$) was verified numerically at $K = 2 \cdot 10^6$, 50 decimal places of precision, on the four (χ, ρ) pairs of § X.5.2. The residuals are:

Pair	residual $\lvert T_K - \mathrm{RHS}(\star)\rvert$ at $K = 2 \cdot 10^6$
χ_{-4}/z_1	$2.85 \cdot 10^{-3}$
χ_{-4}/z_2	$1.66 \cdot 10^{-3}$
χ_5	$4.24 \cdot 10^{-5}$
χ_{11}	$3.33 \cdot 10^{-5}$

These match the predicted $K^{-1/2}/\log K$ decay envelope at the $O(1)$ implicit-constant level (with the χ_{-4} pairs showing the predicted slower envelope due to the $p = 2$ bad-prime weight; see § A.5). The full table with per-component breakdown is at § X.5.4 of the main text.

Appendix B. Appendix B — Full proof of Theorem X.4.2 (c_K leading + subleading identity)

This appendix gives the full proof of Theorem X.4.2 of the main text: under the hypotheses of Theorem X.4.1 (i.e., χ primitive non-principal of conductor q, $\rho = \frac{1}{2} + i\tau$ a *simple* zero of $L(s, \chi)$ with $\tau \neq 0$), and assuming every off-target nontrivial zero $\rho' \neq \rho$ of $L(s, \chi)$ in the truncation height $|\gamma'| \leq T_K$ of the Inoue (2021) explicit-formula box is also *simple*,

$$(\dagger) \qquad c_K(\chi, \rho) \;=\; \frac{\log K}{L'(\rho, \chi)} \;+\; C_1(\chi, \rho) \;+\; o(1) \qquad (K \to \infty),$$

where $C_1(\chi, \rho) = -L''(\rho, \chi)/(2\, L'(\rho, \chi)^2)$ and $c_K(\chi, \rho) = \sum_{n \leq K} \mu(n)\, \chi(n)\, n^{-\rho}$.

The identity (†) is **unconditional in** ρ given simplicity of ρ and of the crossed off-target zeros. The $o(1)$ rate $O(K^{-1/2+\varepsilon})$ for every $\varepsilon > 0$ requires RH for $L(s, \chi)$; the *unconditional* rate (which still yields $o(1)$) is $O(K^{-1/2} \exp((\log K)^{1/2} (\log\log K)^{14}))$ via Soundararajan (2009) Theorem 1.

B.1. B.1 Setup — Perron formula for c_K via $1/L(s, \chi)$.

B.1.1. *B.1.1 The Dirichlet series for* $1/L(s, \chi)$. For χ primitive of conductor $q \geq 2$ and $\mathrm{Re}(s) > 1$,

$$\frac{1}{L(s, \chi)} \;=\; \sum_{n=1}^{\infty} \frac{\mu(n)\, \chi(n)}{n^s},$$

absolutely convergent on $\mathrm{Re}(s) > 1$. The Dirichlet coefficients are $a_n = \mu(n)\, \chi(n)$, and the partial sum $c_K(\rho, \chi) = \sum_{n \leq K} a_n n^{-\rho}$ is the spectroscope object.

B.1.2. *Perron formula (truncated).* The standard truncated Perron formula (Tenenbaum, *Introduction to Analytic and Probabilistic Number Theory*, Theorem II.2.3; Inoue (2021), *Some explicit formulas for partial sums of Möbius functions*, Journal de Théorie des Nombres de Bordeaux **33** (2021), 273–315, eq. (4.1) and Theorem 1) gives the following representation.

For any $\sigma_0' > \frac{1}{2}$ and $T \geq 2$, with the convention that c_K uses the half-weight at integer K,

$$(11) \qquad c_K(\rho, \chi) \;=\; \frac{1}{2\pi i} \int_{\sigma_0' - iT}^{\sigma_0' + iT} \frac{K^w}{w} \cdot \frac{1}{L(w + \rho, \chi)}\, dw \;+\; \mathcal{R}'(K, T, \rho),$$

where the kernel is K^w/w and the Dirichlet series factor is $1/L(w + \rho, \chi)$, obtained by the substitution $s = w + \rho$. The truncation error $\mathcal{R}'(K, T, \rho)$ is bounded by Inoue (2021) Theorem 1 via the J_1, J_2, J_3 estimates of that paper:

$$\mathcal{R}'(K, T, \rho) \;\ll\; \frac{K^{\sigma_0' - 1/2} \log K}{T} \;+\; \min\!\Big(1,\, \frac{K^{1/2}}{T \langle K \rangle}\Big),$$

with $\langle K \rangle$ the distance from K to the nearest integer.

B.1.3. *Inoue (2021) Theorem 1 — truncated explicit formula.* The following statement is verbatim from Inoue (2021), arXiv:1805.05015v1, page 3:

> **Theorem 1** (Inoue 2021). Let $x > 0$, $q \geq 2$, $T \geq \max\{T_0, \exp(q^{1/3}), 2/x\}$. Then, uniformly for all primitive Dirichlet characters χ modulo d with $d \leq q$, there exists $T_\nu \in [T, 2T]$ satisfying

$$M^*(x, \chi) \;=\; \sum_{|\gamma| < T_\nu} \frac{1}{(m(\rho) - 1)!} \lim_{s \to \rho} \frac{d^{m(\rho)-1}}{ds^{m(\rho)-1}} \Big((s - \rho)^{m(\rho)} \frac{x^s}{L(s, \chi)\, s} \Big) \;+\; \operatorname*{Res}_{s=0} \Big(\frac{x^s}{L(s, \chi)\, s} \Big) \;+\; \cdots$$

This is the unshifted form, applied to the partial-character-Möbius sum $M^*(x,\chi) = \sum'_{n\le x} \mu(n)\,\chi(n)$. Dividing by K^ρ and specializing at $x = K$, $m(\rho) = 1$ (the simplicity hypothesis on ρ), and after the change of variable $w = s - \rho$, one obtains (11).

B.2. **B.2 Pole-structure analysis at $w = 0$.** The pole-structure computation of this section is independently formalised in Lean 4 / Mathlib v4.28.0 as `local_perron_residue` in `formal-conjectures/LocalPerr` The Lean theorem states the residue identity as a `Tendsto` limit at the canonical form L analytic with a simple zero at 0 (the general-ρ case of this appendix reduces to the canonical form by the substitution $L \mapsto L(\,\cdot\, + \rho)$); the proof is **unconditional** (0 `sorry`, only `propext`, `Classical.choice`, `Quot.sound` axioms).

We work in the shifted variable $w = s - \rho$. Define

$$F(w) \;:=\; \frac{K^w}{w \cdot L(w+\rho,\chi)}.$$

Since $L(w+\rho,\chi)$ has a *simple* zero at $w = 0$ (by hypothesis), $1/L(w+\rho,\chi)$ has a simple pole at $w = 0$. Combined with the simple pole of K^w/w at $w = 0$, the integrand F has a **double pole at $w = 0$**.

B.2.1. *B.2.1 Laurent expansion of* $1/L(w+\rho,\chi)$. By Taylor expansion at the simple zero,

$$L(w+\rho,\chi) \;=\; a_1 w + a_2 w^2 + a_3 w^3 + \cdots, \qquad a_k \;=\; \frac{L^{(k)}(\rho,\chi)}{k!}, \quad a_1 = L'(\rho,\chi) \neq 0.$$

Therefore

$$\frac{1}{L(w+\rho,\chi)} \;=\; \frac{1}{a_1 w} \cdot \frac{1}{1 + (a_2/a_1)w + (a_3/a_1)w^2 + \cdots}.$$

Expanding the geometric factor as $(1+u)^{-1} = 1 - u + u^2 - \cdots$ and substituting $a_1 = L'(\rho,\chi)$, $a_2 = L''(\rho,\chi)/2$:

$$\frac{1}{L(w+\rho,\chi)} \;=\; \frac{1}{L'(\rho,\chi)\,w} \;-\; \frac{L''(\rho,\chi)}{2\,L'(\rho,\chi)^2} \;+\; O(w). \tag{12}$$

B.2.2. *B.2.2 Laurent expansion of* K^w/w. Trivially,

$$\frac{K^w}{w} \;=\; \frac{1}{w}\big(1 + (\log K)\,w + \tfrac{1}{2}(\log K)^2 w^2 + O(w^3)\big) \;=\; \frac{1}{w} \;+\; \log K \;+\; \tfrac{1}{2}(\log K)^2\,w \;+\; O(w^2).$$

B.2.3. *Product expansion and the double-pole residue.* Multiplying the two expansions:

$$F(w) \;=\; \Big[\tfrac{1}{w} + \log K + \tfrac{1}{2}(\log K)^2 w + \cdots\Big] \cdot \Big[\tfrac{1}{L'(\rho,\chi)\,w} - \tfrac{L''(\rho,\chi)}{2L'(\rho,\chi)^2} + O(w)\Big].$$

Collecting powers of w: - w^{-2} coefficient: $\frac{1}{L'(\rho,\chi)}$ (from $\frac{1}{w}\cdot\frac{1}{L'(\rho,\chi)\,w}$); - w^{-1} coefficient: $\frac{\log K}{L'(\rho,\chi)} - \frac{L''(\rho,\chi)}{2\,L'(\rho,\chi)^2}$ (cross-terms).

Hence

$$\operatorname*{Res}_{w=0} F(w) \;=\; \frac{\log K}{L'(\rho,\chi)} \;-\; \frac{L''(\rho,\chi)}{2\,L'(\rho,\chi)^2} \;=\; \frac{\log K}{L'(\rho,\chi)} + C_1(\chi,\rho). \tag{13}$$

This is **Lemma X.3.1 of the main text**, established algebraically; the local part of (†) is now done. It remains to control the off-target zero residues, the trivial-zero residues, and the contour pieces.

B.3. **Off-target zero residues, trivial-zero residues, and contour pieces.** We now shift the Perron contour from $\mathrm{Re}(w) = \sigma_0' > \frac{1}{2}$ leftward to a zero-avoiding line $\mathrm{Re}(w) = -A$ for some $A > 0$, crossing all the nontrivial zeros of $L(s,\chi)$ in the strip $-A \le \mathrm{Re}(w) \le \sigma_0'$. The integrand $F(w)$ has poles at $w = 0$ (the target double pole, with residue (13)), at $w = \rho' - \rho$ for every other nontrivial zero $\rho' \neq \rho$ of $L(s,\chi)$, at trivial zeros, and at $w = -\rho$ from a (potential) pole of $L(s,\chi)$ at $s = 0$ — but $L(s,\chi)$ is entire for primitive non-principal χ, so the pole at $w = -\rho$ from $1/L$ is absent. We do pick up the term $\mathrm{Res}_{s=0}(x^s/(L(s,\chi)\,s))$ from Inoue's formula (this is the $s = 0$ residue in the unshifted form; after dividing by K^ρ it contributes a term of magnitude $K^{-\rho}/(L(0,\chi)) = O(K^{-1/2})$ in modulus, $o(1)$).

B.3.1. *Off-target nontrivial-zero residues.* For each $\rho' \neq \rho$ a nontrivial zero of $L(s,\chi)$, write $w_{\rho'} = \rho' - \rho$. Assuming ρ' is *simple* (which is generically expected and is the regime relevant for the leading + subleading identity), the residue of F at $w = w_{\rho'}$ is

$$\operatorname*{Res}_{w=w_{\rho'}} F(w) \;=\; \frac{K^{w_{\rho'}}}{w_{\rho'} \cdot L'(\rho',\chi)} \;=\; \frac{K^{\rho'-\rho}}{(\rho'-\rho)\,L'(\rho',\chi)}.$$

Under RH for $L(s,\chi)$, $|K^{\rho'-\rho}| = 1$ for all $\rho' = \frac{1}{2} + i\gamma'$ with $\gamma' \in \mathbb{R}$. The aggregate off-target sum (over zeros ρ' in the strip $|\gamma'| \le T$) is then bounded by

$$\Big|\sum_{\rho'\neq\rho,\ |\gamma'|\le T} \frac{K^{i(\gamma'-\tau)}}{(\rho'-\rho)\,L'(\rho',\chi)}\Big| \;\ll\; \sum_{\rho'\neq\rho,\ |\gamma'|\le T} \frac{1}{|\gamma'-\tau|\,|L'(\rho',\chi)|}.$$

This is the off-target *aggregate* whose control determines the $o(1)$ rate.

B.3.2. *Trivial-zero residues.* Trivial zeros of $L(s,\chi)$ are at $s = -2k$ (χ even) or $s = -2k-1$ (χ odd), $k \ge 0$. After the shift $w = s - \rho$, these contribute residues of size $K^{-\mathrm{Re}(\rho)-2k} = K^{-1/2-2k}$, summable to $O(K^{-1/2})$, which is $o(1)$.

B.3.3. *Shifted contour pieces.* The shifted contour $\mathrm{Re}(w) = -A$ contributes, on a suitable zero-avoiding height $T_K \in [K/(\log K)^B, 2K/(\log K)^B]$, an integral bounded by (using a reciprocal-L bound on the vertical shifted line)

$$\Big|\frac{1}{2\pi i}\int_{(-A)} \frac{K^w}{w\,L(w+\rho,\chi)}\,dw\Big| \;\ll\; K^{-A} \cdot T_K \cdot \exp\big(C\,(\log\log T_K)^2\big).$$

For $A > 1/2$ and $T_K \le K/(\log K)^B$ with B sufficiently large, this is $K^{-A+1+o(1)} \cdot \exp(O((\log\log K)^2))$, which is $o(1)$.

Horizontal pieces give factors like $(T_K/K)^A \cdot$polylog, $o(1)$ for $T_K \le K(\log K)^{-B}$ and large enough B.

B.3.4. *The Perron truncation error.* By Inoue (2021) Theorem 1, the truncation error $\mathcal{R}'(K, T_K, \rho)$ is bounded by

$$\mathcal{R}'(K, T_K, \rho) \;\ll\; \frac{K^{\sigma_0'-1/2}\log K}{T_K} \;+\; \min\Big(1,\, \frac{K^{1/2}}{T_K\langle K\rangle}\Big).$$

For $T_K \asymp K(\log K)^{-B}$ and $\sigma_0' = 1/2 + 1/\log K$, both pieces are $O(K^{-1/2+\varepsilon}) \cdot (\log K)^{O(1)}$ for any $\varepsilon > 0$, in particular $o(1)$.

B.3.5. *Assembly.* Combining the local residue (13) with the off-target and contour pieces:

$$c_K(\chi,\rho) \;=\; \frac{\log K}{L'(\rho,\chi)} + C_1(\chi,\rho) + \Big(\text{off-target aggregate}\Big) + O(K^{-1/2}) + O(K^{-A+1+o(1)}\exp(O((\log\log K)^2))) + \mathcal{R}$$

All terms in the last three parentheses are $o(1)$. The off-target aggregate is also $o(1)$ under RH for $L(s,\chi)$, by partial summation against Soundararajan (2009), *Partial sums of the Möbius function*, Annals of Mathematics **170** (2009), 981–993, Theorem 1:

$$\text{(Soundararajan 2009)}\quad \text{under RH:}\quad M(x) \;\ll\; \sqrt{x}\,\exp\!\big((\log x)^{1/2}\,(\log\log x)^{14}\big).$$

Translating to $M^*(x,\chi)$ (the character analogue) gives the required $o(1)$ rate for the off-target aggregate under RH for $L(s,\chi)$.

This establishes (†). □

B.4. **B.4 Rate of the $o(1)$.** The proof above yields the following rates.

Rate	Hypotheses
$o(1) = O\big(K^{-1/2+\varepsilon}\big)$ for every $\varepsilon > 0$	RH for $L(s,\chi)$.
$o(1) = O\big(K^{-1/2}\exp\big((\log K)^{1/2}(\log\log K)^{14}\big)\big)$	**unconditional** (Soundararajan 2009 applies under RH for $L(s,\chi)$, which holds for the four characters $\chi_{-4},\chi_5,\chi_{11},\ldots$ at the verified heights).
$o(1) = O\big(\log K/\sqrt{K}\big)$	RH plus a Gonek–Hejhal-type bound of the form $\sum_{0<\gamma<T} \lvert L'(\rho,\chi)\rvert^{-2} \ll T(\log T)^{O(1)}$.

In the numerical work of § X.5.3, the empirical residual $R(K) := c_K - \log K/L'(\rho,\chi) - C_1(\chi,\rho)$ at $K = 2\cdot 10^6$ falls in the interval [0.027, 0.375] across the four pairs; this is consistent with the Soundararajan-conditional rate within an $O(1)$ implicit-constant factor (specifically, $\lvert R(K)\rvert/(\log K/\sqrt{K}) \in$ [3, 36], well within the typical $\exp(C(\log\log K)^{1/2}(\log\log K)^{14})$ envelope at this scale).

B.5. **B.5 What the proof does *not* claim.** For clarity:

- The proof does **not** establish the Aoki–Koyama–Mertens limit for $E_K \log K$, which is the *multiplicative-side* asymptotic.
- The proof does **not** establish the *full* Perron-leading theorem $c_K = \log K/L'(\rho) + o(\log K)$ unconditionally; the off-target aggregate's $o(\log K)$ control (which is *weaker* than the $o(1)$ obtained in (†) at the *leading + subleading* level) depends on RH for $L(s,\chi)$ and is what is required for composing with (AK) to obtain (NDC).
- The proof does **not** address the possibility of off-target *multiple* zeros, which under DRH/EDRH alone are not excluded; the manuscript's open Question Q:Perron asks for either the off-target multiple-zero residue control or a sufficient "all crossed off-target zeros simple" hypothesis.